

mmWave UE Positioning Under Blockage: Limits for ISAC-Assisted Handover

Seyed Bagher Hashemi Natanzi^{*}, Bo Tang[†]

^{*†}Department of Electrical and Computer Engineering, Worcester Polytechnic Institute (WPI), USA

Email: {snatanzi, btang1}@wpi.edu

Abstract—Blockage-aware proactive handover in millimeter-wave (mmWave) networks predicts when a body will cross the line of sight (LoS) between a radio unit and a user, and is therefore only as good as the network’s knowledge of the user position. This paper studies whether the radio units that serve a user can localize it accurately enough from its uplink sounding reference signals. We derive single-epoch position error bounds that include synchronization and array calibration errors as Bayesian priors, contrast LoS-only information with map-aided information from a digital twin, and evaluate a practical estimator with an extended Kalman filter. In ray-traced experiments in a street canyon, reported on held-out scenarios, the bound is limited by array calibration rather than bandwidth, blockage substantially degrades the LoS-only bound, and known reflected paths preserve a much lower map-aided bound. Fixed-geometry experiments reveal substantial systematic errors in the evaluated estimator, alongside geometry-dependent random error, and its errors grow sharply around blockage events. Fed into a blockage-aware handover planner, its positions do not let foresight outperform 3GPP event-triggered handover, whereas synthetic errors at the level of the map-aided bound do.

Index Terms—mmWave positioning, position error bound, uplink SRS, blockage, proactive handover, digital twin, ISAC.

I. INTRODUCTION

Millimeter-wave (mmWave) bands provide the bandwidth that 5G and 6G networks need for high data rates, but their links are fragile: a pedestrian or vehicle crossing the line of sight (LoS) attenuates the signal by tens of decibels within milliseconds [1], [2]. Reactive handover cannot avoid such outages. Integrated sensing and communication (ISAC) and other sensing-aided approaches therefore aim to predict blockage and hand the user over in advance [3]–[5].

Predicting a blockage means predicting whether a body will cross the line between a radio unit and the user. The prediction can thus be no better than the network’s knowledge of where the user is. In a companion study [6], we found that a blockage-aware handover planner outperforms the service-aware 3GPP A5 event only if the user position is known with an error of about 8–13 cm (per horizontal axis), and that a meter-level error removes the benefit entirely. Positioning accuracy is therefore not a side issue but a prerequisite of sensing-aided mobility.

Can the network meet this requirement with the infrastructure it already has? 5G NR supports uplink positioning from sounding reference signals (SRS) [7], and recent street-scale

studies with several base stations report errors below a meter [8]. Position error bounds for mmWave systems suggest that much better accuracy is possible in principle [9], [10], but they usually assume ideal hardware and static scenes. It remains unclear which effect limits decimeter accuracy in practice, how blockage changes it, how close practical estimators come, and whether their errors are small enough at the moments when a handover decision is made.

This paper addresses these questions in a street canyon served by two O-RAN radio units (O-RUs). Our contributions are:

- A positioning analysis tailored to blockage-aware handover: a single-epoch equivalent Fisher information analysis with synchronization and array calibration errors as Bayesian priors, two models of a blocked LoS, and map-aided information from reflected paths with known geometry.
- A practical uplink estimator with an extended Kalman filter (EKF), compared with the bound at fixed geometries by decomposing its error into bias and random spread.
- An analysis of positioning errors around blockage onsets and a closing experiment that feeds estimated and synthetic bound-level positions into a blockage-aware handover planner. The bounds, the estimator, and the scripts that generate every number in this paper are released as open source.¹
- Findings on held-out scenarios: calibration rather than bandwidth limits the bound; blockage substantially degrades the LoS-only bound, while known reflections preserve a much lower map-aided bound; practical errors are mostly systematic at the tested geometries; and they worsen in the time windows relevant to proactive handover.

II. RELATED WORK

Bounds for mmWave positioning. Fisher information analyses show that the angle and delay resolution of mmWave arrays support sub-decimeter accuracy [9], [10], building on the equivalent Fisher information framework for wideband localization [11]. Multipath components with known geometry carry position information as virtual anchors [12], [13]. We use these tools, add hardware errors as priors, and evaluate them under traffic-driven blockage.

¹<https://github.com/CLIS-WPI/sense2connect> (branch `paper2`, tag `v1.0-wcnc2027-paper2`).

^{*}Corresponding author: snatanzi@wpi.edu

Uplink SRS positioning. 3GPP specifies uplink time- and angle-based positioning [7]. Dong *et al.* [8] combine LoS/NLoS gating, dominant-path angle and delay estimation, and multi-base-station EKF fusion, reach sub-meter errors in a ray-traced campus, and point to reflections as virtual anchors for future work. Our estimator follows a similar architecture; our question is whether its accuracy suffices for a specific downstream use.

Sensing-aided handover. Vision-, radar-, and ISAC-aided handover predicts blockage or user motion [4], [5]; our companion study [6] identified the user position as its binding requirement.

III. SYSTEM MODEL AND PROBLEM FORMULATION

A. Signal Model

A single-antenna user equipment (UE) at horizontal position $\mathbf{p}$ transmits an uplink SRS that is received by a set $\mathcal{C}$ of O-RUs, each with an M -element array. At O-RU c , subcarrier frequency f_n , and element m , the received sample is

$$y_{c,n,m} = \sum_{\ell} \beta_{c,\ell} e^{-j2\pi f_n(\tau_{c,\ell} + \delta_c + b)} e^{j\psi_{c,m}} s_m(\mathbf{u}_{c,\ell}) + w_{c,n,m}, \quad (1)$$

where path ℓ has complex gain $\beta_{c,\ell}$ (including blockage loss), delay $\tau_{c,\ell}$, and arrival direction $\mathbf{u}_{c,\ell}$ with steering response s_m ; δ_c is the time offset of O-RU c , b the UE clock bias, $\psi_{c,m}$ the calibration phase error of element m , and $w_{c,n,m}$ thermal noise. Hardware errors are random with $\delta_c \sim \mathcal{N}(0, \sigma_{\text{sync}}^2)$ and $\psi_{c,m} \sim \mathcal{N}(0, \sigma_{\phi}^2)$.

B. Geometry and Blockage

The LoS delay and angles are functions of $\mathbf{p}$ and the O-RU position $\mathbf{p}_c$, e.g., $\tau_{c,0} = \|\mathbf{p} - \mathbf{p}_c\|/c_0$. A specular reflection on a known surface is an LoS path from the virtual anchor $\tilde{\mathbf{p}}_{c,\ell}$, the image of $\mathbf{p}_c$ in that surface, so that $\tau_{c,\ell} = \|\mathbf{p} - \tilde{\mathbf{p}}_{c,\ell}\|/c_0$ [12]. Moving bodies attenuate every path segment according to 3GPP blockage model B [14]. An O-RU is *blocked* for a UE at an epoch when the model-B loss of its LoS path is at least 10 dB. Model B reduces the power of a blocked LoS but keeps its geometry, whereas the field that actually arrives diffracts around the blocker. We therefore consider two models of a blocked LoS: *kept*, in which its delay and angles still follow $\mathbf{p}$ (optimistic), and *biased*, in which they carry an unknown bias. In the biased model, the bound treats them as free parameters, and the signal fed to the estimator replaces the blocked LoS by a detour around its dominant blocker (the blocker with the largest model-B loss): with $\mathbf{q}$ the shorter of the detour points over the top and past either side of the blocker box, next to the point of the LoS closest to the blocker, the path keeps the attenuated complex gain of the LoS but has delay $(\|\mathbf{q} - \mathbf{p}_c\| + \|\mathbf{p} - \mathbf{q}\|)/c_0$ and arrival direction along $\mathbf{q} - \mathbf{p}_c$. Reflected paths are only attenuated.

C. Positioning Requirement of Blockage-Aware Handover

The handover planner of [6] receives tracks of the blockers and the reported UE position every 0.1 s, predicts the signal-to-noise ratio (SNR) of each cell by applying model B on

the line between each O-RU and the reported position, and chooses the serving cell sequence that minimizes the predicted outage over a horizon of 0.5–3 s, accounting for an interruption of 20 ms per handover; decisions take effect after a control-loop delay of 20 ms, and the sensing overhead of 2.3% of the resources is charged to the planner. Its performance is measured by the outage time at a service rate, compared with the 3GPP A5 event. With perfect blocker tracks and a white Gaussian UE position error, the planner beats A5 only up to a per-axis standard deviation of about 8–13 cm, i.e., an RMS distance error of about 11–18 cm or a median distance error of about 9–15 cm. The positioning problem is thus to estimate $\mathbf{p}$ at every decision epoch such that the resulting handover outage stays below that of A5; we examine (i) what limits the bound, (ii) how blockage and map knowledge change it, (iii) how a practical estimator compares, and (iv) whether its positions suffice for the planner.

IV. POSITION ERROR BOUND AND ESTIMATOR

A. Single-Epoch Position Error Bound

Collect the channel parameters of all O-RUs at one epoch in $\boldsymbol{\eta}$ (delays, angles, complex gains) and the hardware and timing errors in $\boldsymbol{\nu} = [\{\delta_c\}, b, \{\psi_{c,m}\}]$. The Fisher information of $\boldsymbol{\eta}$ follows from (1) by the Slepian–Bangs formula; we compute it exactly by factorizing its frequency and antenna sums. With $\mathbf{T} = \partial\boldsymbol{\eta}/\partial\mathbf{p}$, the equivalent Fisher information of the position is

$$\mathbf{J}_{\mathbf{p}} = \mathbf{T}^T \mathbf{J}_{\boldsymbol{\eta}} \mathbf{T} - \mathbf{T}^T \mathbf{J}_{\boldsymbol{\nu}\boldsymbol{\nu}} (\mathbf{J}_{\boldsymbol{\nu}\boldsymbol{\nu}} + \boldsymbol{\Lambda})^{-1} \mathbf{J}_{\boldsymbol{\nu}\boldsymbol{\eta}} \mathbf{T}, \quad (2)$$

after the nuisance parameters have been removed by Schur complements: the complex gains of all paths; in the *LoS-only* bound, the delays and angles of all reflected paths; in the biased model, the delay and angles of a blocked LoS; and the timing and hardware errors in $\boldsymbol{\nu}$. $\boldsymbol{\Lambda}$ holds the Gaussian priors of the O-RU time offsets (σ_{sync}) and of the element phases (σ_{ϕ}); the UE clock bias has no prior (time difference of arrival, TDoA). In the *map-aided* bound, the reflected paths are not nuisance parameters: their geometry is known from the digital twin, and their entries of $\mathbf{T}$ follow analytically from the virtual anchors. In double precision, the analytic derivatives of delay and angles agree with central finite differences for single- and double-bounce paths at 1000 random UE positions to a relative error of at most 3.2×10^{-6} (step 10^{-4} m), with the expected $O(h^2)$ behavior for larger steps. Against single-precision re-tracing with Sionna at ± 1 cm, 77% of the derivative values agree to a relative error of 10^{-3} ; of the disagreeing paths, 86% are explained by single-precision rounding (absolute error growing as $1/h$) and 14% by derivatives that are exactly zero, with 0% unexplained. The GPU implementation of (2) matches an independent NumPy implementation to a relative deviation of 2.9×10^{-6} and a brute-force numerical Fisher matrix on toy cases to 6.6×10^{-10} . The position error bound is $\text{PEB} = \sqrt{\text{tr}(\mathbf{J}_{\mathbf{p}}^{-1})}$. It bounds the root-mean-square (RMS) error of unbiased single-epoch estimates for a given geometry; it is neither a median error nor a guarantee that an estimator attains it, a biased estimator can have a smaller spread, and




Fig. 1. Scenario geometry (dimensions indicated; axes scaled independently). (a) Two neighboring cells share the street; O-RU 0 and O-RU 1 are on 5 m lampposts, 40 m apart along the street. The true serving link (solid red) crosses the bus, whereas the link predicted from the reported UE position (dotted) passes it, so a position error turns into a missed blockage prediction; green: link to the other cell and one of its facade reflections (dashed). (b) The LoS passes through a bus but above a car. (c) In the facade deployment, O-RU 0 is mounted on the north facade at 8 m.

it excludes the temporal information that a tracking filter accumulates.

B. Practical Estimator

Per O-RU and epoch, a zero-padded two-dimensional DFT over the array selects the strongest beam; a Hann-windowed delay transform of the beamformed signal gives the delay of the dominant path, refined by parabolic interpolation and followed by an angle refinement at that delay. A peak-to-second-peak test flags reflected dominant paths. Because a planar array of isotropic elements cannot distinguish the two sides of its plane, each angle yields two direction hypotheses; the estimator keeps the joint hypothesis that is consistent across O-RUs, lies in the walkable area of the digital-twin map, and passes the gate of a constant-velocity EKF, which fuses weighted least-squares fixes over time.

V. EXPERIMENTS AND RESULTS

A. Setup

Scenario. We use the street canyon of [6] (Fig. 1): buildings 18.2 m apart, two 3.5-m lanes, and sidewalks on both sides; actors and UEs move on a periodic section of 80 m. Two O-RUs with 8×8 half-wavelength arrays at 28 GHz form two neighboring cells that share this section: O-RU 0 on the north side, either on a lamppost (5 m high) or on a facade (8 m), and O-RU 1 on a 5-m lamppost on the south sidewalk, 40 m further along the street. Pedestrian UEs at 1.5 m walk along the south sidewalk at 1.0–1.2 m/s. Each 60-s run has low or high traffic: 37.5 or 100 vehicles per lane and km (cars, buses, trucks at 6–12 m/s) and 6 or 12 pedestrians on the sidewalks plus 2 or 4 crossing the street. Paths are traced with Sionna RT [15], [16] (specular, order up to 2, at most 8 paths per link).

Service model. The SNR of a cell adds, over all traced paths, the path power with the full array gain after its model-B loss; this is an optimistic approximation of the received multipath power, since no single beam collects all paths with full gain. The UE is in outage when this total remaining power gives an SNR below $\gamma_{\text{req}} = 9.3$ dB, the value needed for 400 Mbit/s. A LoS blockage thus causes an outage only if the remaining paths do not sustain γ_{req} : during 10 dB LoS events, the strongest surviving path of the same cell lies a median 26–33 dB below the unblocked LoS [6], so the link margin determines whether the same cell can still serve the UE or the other cell is needed.

Signals and hardware. The UE transmits one SRS symbol per 0.1-s epoch at 23 dBm; the O-RU noise figure is 7 dB. Numerology 3 (120 kHz) with 100, 200, and 400 MHz carriers uses 792, 1584, and 3168 active subcarriers [17]. We sweep σ_{sync} of 0, 0.3, 1 and 3 ns (drawn once per run, a bias that the EKF cannot average out) and σ_{ϕ} of 0, 2 and 5 deg (fixed per O-RU and run). The main configuration uses TDoA, 400 MHz, $\sigma_{\text{sync}} = 1$ ns, $\sigma_{\phi} = 2$ deg, and the biased blocked-LoS model.

Protocol. Estimator parameters were tuned on seeds 101–105 and methods fixed on development seeds 1001–1010. An earlier held-out set was consumed before the front/back disambiguation was added, so all results come from a second, never-used set of ten seeds (3001–3010), each simulated for two mounts and two traffic densities. The seed is the statistical unit: we average its runs, compare conditions with the exact two-sided Wilcoxon signed-rank test over the ten seeds, and give 95% bootstrap intervals over seeds. All reported numbers, tables, and figures are generated by the released scripts¹ from the stored results, including the fixed-geometry validation and the derivative checks of Section IV.

TABLE I

SINGLE-EPOCH ESTIMATOR ERROR AT FIXED GEOMETRIES (500 DRAWS EACH, MAIN HARDWARE, TDOA, NO EKF) VERSUS THE SINGLE-EPOCH LoS-ONLY PEB (BIASED MODEL FOR A BLOCKED LoS); ALL VALUES IN CM; $\text{RMSE}^2 = \text{BIAS}^2 + \text{CENTERED RMS}^2$ (EQ. (3)).

Mount	LoS state	UE x [m]	PEB	400 MHz			100 MHz			
				bias	centered	RMSE	PEB	bias	centered	RMSE
lamppost	both LoS	-20.0	1.5	1.3	1.7	2.2	1.5	7.7	7.4	10.7
lamppost	both LoS	20.0	4.2	1496.1	55.4	1497.1	4.5	875.6	345.0	941.2
lamppost	blocked (O-RU 0)	21.9	46.3	301.3	4.9	301.3	91.5	286.1	9.3	286.3
facade	both LoS	-20.0	3.3	31.9	4.8	32.2	3.4	30.0	16.1	34.1
facade	both LoS	20.0	6.7	88.6	6.1	88.8	7.1	94.0	9.4	94.5
facade	blocked (O-RU 0)	21.9	90.8	189.2	7.0	189.3	208.0	185.0	10.4	185.3




Fig. 2. Median single-epoch position error bound (LoS-only, map-aided) and median estimator error versus bandwidth, for epochs with LoS to both O-RUs and with a blocked O-RU; main configuration (95% CI). The bound and the error are different quantities (Section IV).

B. What Limits the Bound

Fig. 2 shows the median bounds versus bandwidth. With LoS to both O-RUs, the LoS-only PEB is below 10 cm in 84% of the epochs at 400 MHz and still in 81% at 100 MHz. Relaxing one limit at a time at 200 MHz scales the median LoS-only PEB (4.8 cm) by 0.97 for twice the bandwidth, 0.94 for ten times the SNR, 0.72 for perfect synchronization, and 0.07 for perfect calibration. On these short links, the delay and angle measurements are precise enough that the element phase errors, which bias the angles, dominate. Without hardware errors, the LoS-only PEB is below 10 cm in 92% of the epochs; with 3 ns and 5 deg errors, in 44%. The map-aided PEB is far less sensitive (99% and 98%), because reflections from known surfaces add independent geometric information.

C. Blockage and the Digital Twin

In the biased model, the median LoS-only PEB rises by a factor of 2.8 during blockage ($p = 0.002$), and only 47% of the blocked epochs remain below 10 cm. If the attenuated LoS




Fig. 3. Median bound and estimator error along the sidewalk (5-m bins of the UE position); main configuration.

kept its geometry, the factor would be only 1.2, which shows how optimistic it is to apply model B's power loss without its geometric consequence. Map-aided information lowers the PEB during blockage to 0.09 times the LoS-only value ($p = 0.002$). Along the street (Fig. 3), the LoS-only PEB varies with the geometry of the two O-RUs from about a centimeter to over a decimeter, while the map-aided PEB stays mostly in the millimeter range. A low map-aided PEB indicates that the information is present in the signal; whether an estimator can extract it is a separate question.

D. Practical Estimator and the Bound

At 400 MHz, the estimator reaches a median error of 12 cm (10 cm with LoS to both O-RUs, 27 cm during blockage); 45% of its errors are below 10 cm, and its 90th percentile is 2.10 m. At 100 MHz, the median grows to 24 cm.

To compare the estimator with the bound on equal terms, we ran 500 independent noise and hardware draws through the single-epoch estimator (without EKF) at 6 fixed geometries and decomposed its error as

$$\text{RMSE}^2 = \|\mathbb{E}[\hat{\mathbf{p}}] - \mathbf{p}\|^2 + \mathbb{E}[\|\hat{\mathbf{p}} - \mathbb{E}[\hat{\mathbf{p}}]\|^2], \quad (3)$$

i.e., squared bias plus squared centered RMS (Table I). The squared bias accounts for most of the mean squared error in




Fig. 4. Outage at 400 Mbit/s versus link margin. Except for A5, every curve is the blockage-aware planner with perfect blocker tracks, fed with the indicated UE position (bound-level errors are synthetic).

TABLE II
ESTIMATOR ERROR AROUND REFERENCE-CELL BLOCKAGE EVENTS
(10 dB, MAIN CONFIGURATION).

Window	Median	90th perc.	>0.1 m	>0.25 m
1.0–0.5 s before onset	12 cm	4.24 m	54%	30%
0.5–0 s before onset	11 cm	4.49 m	53%	29%
During the event	87 cm	3.26 m	88%	72%
Elsewhere	10 cm	1.59 m	49%	29%

five of the six geometries at 400 MHz (98–100% of it); in the remaining one (lamppost, both LoS, $x = -20$ m), the squared bias is 38% of it and the random spread is larger. The bias is caused by interference of unresolved multipath, and in one geometry by a front/back mirror solution. The centered RMS is geometry dependent, between 0.077–13 times the PEB at 400 MHz: close to the bound at some geometries and an order of magnitude above it at others. In blocked geometries it falls below the PEB, which is possible because the estimator is biased: it uses the detoured LoS precisely but in the wrong place. Over the held-out runs, front/back mirror solutions remain in 5% of the epochs and dominate the squared error (RMS 6.36 m), while reflected dominant paths set the 90th percentile. A variant that used reflections as virtual anchors did not reduce the tail on the development seeds (90th percentile 3.52 vs. 1.87 m) and was not pursued.

E. Positioning Errors Around Blockage Onsets

Table II shows the estimator error around the 752 blockage events of the reference cell, defined per UE as the cell with the larger unblocked received power; an event is an interval with LoS loss of at least 10 dB, with gaps below 0.5 s merged. Each epoch is counted in exactly one window, with priority to the event, then to the last half second before onset, then to the half

TABLE III
PLANNER OUTAGE MINUS A5 [s/UE-MIN] BY SOURCE OF UE POSITION
(PERFECT BLOCKER TRACKS; COLUMNS: LINK MARGIN, REF. = 3GPP
REFERENCE BUDGET WITH A MARGIN OF 35.9 dB; †: NOT SIGNIFICANT,
SEED-LEVEL WILCOXON $p \geq 0.05$; POINTWISE TESTS, EXPLORATORY).

UE position	15 dB	20 dB	25 dB	30 dB	Ref.
Exact	-0.108	-0.066†	-0.061	-0.007†	-0.025
<i>Synthetic, bound level (full covariance)</i>					
map-aided	-0.099	-0.049†	-0.044	+0.021†	-0.005†
LoS-only	+0.162	+0.301	+0.175	+0.151	+0.097
<i>Estimated</i>					
estimator	+0.652	+0.507	+0.284	+0.118	+0.052
ideal hardware	+0.584	+0.483	+0.295	+0.121	+0.054
capped at p90	+0.525	+0.434	+0.251	+0.118	+0.051
mirror-corr.*	+0.461	+0.443	+0.260	+0.115	+0.048
White 1 m error	+1.402	+0.824	+0.476	+0.290	+0.117

*Oracle diagnostic: mirror solutions replaced using the true position.

second before that; the first 2 s of each run are excluded. In the second before onset, a window relevant to proactive handover, the median error and the share of errors above 0.25 m are about as elsewhere, but the large errors become larger: the 90th percentile rises from 1.59 to 4.49 m. During the event, most of the distribution deteriorates, with the median growing to 87 cm. Errors thus worsen in the time windows relevant to proactive handover: before a blockage, the extreme errors grow; during it, the bulk of the errors does, which the overall median hides. The windows refer to the blockage events of a reference cell, not to the decision times of a particular planner.

F. Implications for Handover

Fig. 4 and Table III feed positions into the planner, with perfect blocker tracks so that the user position is the only source of error. With exact positions, the planner has a lower outage than A5 from 15 dB margin upward (significantly at 15 and 25 dB and at the reference). As a sensitivity test, we feed synthetic errors $\mathbf{e}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{J}_p(t)^{-1})$, drawn independently per epoch from the single-epoch bound of the respective model, which is anisotropic (median ratio of principal standard deviations 3.0). Errors at the map-aided bound preserve a lower mean outage than A5 at most tested margins, with significant improvements at 15 and 25 dB; errors at the LoS-only bound, which grow during blockage, make it worse than A5 at every margin. With the positions of the evaluated estimator, the planner is significantly worse than A5 at every margin, also with ideal hardware. Capping the errors at their 90th percentile or removing the mirror solutions with the true position reduces the loss but leaves it significant, so the largest errors are not its only cause. Synthetic bound-level errors are an optimistic model; they show what positioning would have to achieve, not that a practical estimator achieves it. The absolute gains at stake are small: at 15 dB margin, A5 has an outage of 1.88 s/UE-min, and the relative change is -5.7% with exact positions and -5.2% with bound-level map-aided positions; at the reference, A5 has 0.058 s/UE-min, and the changes are -43.1% and -8.2%†, respectively. These reductions may benefit interruption-sensitive services, although

the application-level impact depends on the frequency and duration of individual outages, which are not evaluated here.

VI. DISCUSSION AND LIMITATIONS

For blockage-aware handover, the relevant accuracy is the error around blockage onsets, not the overall median, and at most tested geometries the gap between the bound and the evaluated estimator is mainly systematic multipath error; these results do not rule out stronger practical estimators. The information to close this gap is present in the reflected paths, as the map-aided bound shows, but our simple virtual-anchor variant did not extract it; map-aided estimators that reliably associate reflections with known surfaces [8], [13] and better array calibration are the main levers. The study has limitations: a single street canyon with two O-RUs; specular paths of order at most 2 and no diffuse scattering; a perfectly known map in the map-aided bound; the biased model of a blocked LoS as our modelling choice; single-epoch bounds without the temporal information of tracking; and planar arrays of isotropic elements, whose front/back ambiguity a physical back plane would remove in hardware [8].

VII. CONCLUSION

We examined whether the mmWave radio units that serve a user can localize it accurately enough for blockage-aware handover. Under the tested configurations, the single-epoch bound is limited by array calibration rather than bandwidth; blockage substantially degrades the LoS-only bound, while known reflected paths preserve a much lower map-aided bound. The evaluated estimator shows substantial systematic multipath errors, and its errors worsen in the time windows relevant to proactive handover; its positions do not let foresight outperform reactive handover. Map-aided estimation that turns reflections into reliable position information is therefore the next step toward ISAC-assisted mobility.

ACKNOWLEDGMENT

This material is based upon work supported in part by NSF under Awards CNS-2120442 and IIS-2325863, and NTIA under Award No. 51-60-IF007. Any opinions, findings, and conclusions or recommendations expressed in this publication are those of the author(s) and do not necessarily reflect the views of the NSF and NTIA.

REFERENCES

- [1] T. S. Rappaport, S. Sun, R. Mayzus, H. Zhao, Y. Azar, K. Wang, G. N. Wong, J. K. Schulz, M. Samimi, and F. Gutierrez, "Millimeter wave mobile communications for 5G cellular: It will work!" *IEEE Access*, vol. 1, pp. 335–349, 2013.
- [2] G. R. MacCartney, T. S. Rappaport, and S. Rangan, "Rapid fading due to human blockage in pedestrian crowds at 5G millimeter-wave frequencies," in *Proc. IEEE GLOBECOM*, 2017.
- [3] F. Liu, Y. Cui, C. Masouros, J. Xu, T. X. Han, Y. C. Eldar, and S. Buzzi, "Integrated sensing and communications: Toward dual-functional wireless networks for 6G and beyond," *IEEE J. Sel. Areas Commun.*, vol. 40, no. 6, pp. 1728–1767, 2022.
- [4] G. Charan, M. Alrabeiah, and A. Alkhateeb, "Vision-aided 6G wireless communications: Blockage prediction and proactive handoff," *IEEE Trans. Veh. Technol.*, vol. 70, no. 10, pp. 10 193–10 208, 2021.
- [5] Y. Li, F. Liu, H. Xiong, Z. Wang, Narengerile, and C. Masouros, "Look before switch: Sensing-assisted handover in 5G NR V2I networks," *IEEE Trans. Mobile Comput.*, 2026, early access, doi: 10.1109/TMC.2025.3647295.
- [6] S. B. Hashemi Natanzi and B. Tang, "ISAC-assisted proactive handover in mmWave networks: Value of foresight and localization requirements," submitted to IEEE WCNC, 2027.
- [7] 3GPP, "Study on NR positioning support," 3rd Generation Partnership Project, Technical Report TR 38.855, Release 16, 2019.
- [8] W. Dong, Z. Ju, H. Jiao, R. Chen, and L. Chen, "Deployable uplink SRS-based 5G NR positioning in mixed LoS/NLoS environments," *Expert Syst. Appl.*, vol. 328, p. 132887, 2026, doi: 10.1016/j.eswa.2026.132887.
- [9] A. Shahmansoori, G. E. Garcia, G. Destino, G. Seco-Granados, and H. Wymeersch, "Position and orientation estimation through millimeter-wave MIMO in 5G systems," *IEEE Trans. Wireless Commun.*, vol. 17, no. 3, pp. 1822–1835, 2018.
- [10] Z. Abu-Shaban, X. Zhou, T. Abhayapala, G. Seco-Granados, and H. Wymeersch, "Error bounds for uplink and downlink 3D localization in 5G millimeter wave systems," *IEEE Trans. Wireless Commun.*, vol. 17, no. 8, pp. 4939–4954, 2018.
- [11] Y. Shen and M. Z. Win, "Fundamental limits of wideband localization—Part I: A general framework," *IEEE Trans. Inf. Theory*, vol. 56, no. 10, pp. 4956–4980, 2010.
- [12] E. Leitinger, P. Meissner, C. Rüdiger, G. Dumphart, and K. Witrisal, "Evaluation of position-related information in multipath components for indoor positioning," *IEEE J. Sel. Areas Commun.*, vol. 33, no. 11, pp. 2313–2328, 2015.
- [13] K. Witrisal, P. Meissner, E. Leitinger, Y. Shen, C. Gustafson, F. Tufveson, K. Haneda, D. Dardari, A. F. Molisch, A. Conti, and M. Z. Win, "High-accuracy localization for assisted living: 5G systems will turn multipath channels from foe to friend," *IEEE Signal Process. Mag.*, vol. 33, no. 2, pp. 59–70, 2016.
- [14] 3GPP, "Study on channel model for frequencies from 0.5 to 100 GHz," 3rd Generation Partnership Project, Technical Report TR 38.901, Release 19, 2025.
- [15] J. Hoydis, F. Ait Aoudia, S. Cammerer, M. Nimier-David, N. Binder, G. Marcus, and A. Keller, "Sionna RT: Differentiable ray tracing for radio propagation modeling," in *Proc. IEEE GLOBECOM Workshops*, 2023.
- [16] J. Hoydis, S. Cammerer, F. Ait Aoudia, M. Nimier-David, L. Maggi, G. Marcus, A. Vem, and A. Keller, "Sionna, version 2.2.0," <https://nvlabs.github.io/sionna/>, 2022.
- [17] 3GPP, "NR; user equipment (UE) radio transmission and reception; part 2: Range 2 standalone," 3rd Generation Partnership Project, Technical Specification TS 38.101-2, 2024.